

NLP

from
Scratch


TEXT

UNDERSTAND

GENERATE

Understand
Language

Generate
Meaningful
Responses



✦ NATURAL LANGUAGE PROCESSING (NLP) ✦



NLP is a field of AI that helps computers **understand, interpret** and **generate** human language.

GOAL:

To bridge the gap between Human Communication and Machine Understanding.

1) CORE NLP TASKS ☆

- Text Classification (e.g. Spam Detection)
- Sentiment Analysis (e.g. Positive/Negative)
- Named Entity Recognition (NER)
- Machine Translation
- Text Summarization
- Question Answering
- Language Generation
- Chatbots & Dialog Systems



2) NLP WORKFLOW ☆

- 1) Text Input
- 2) Preprocessing
- 3) Feature Representation
- 4) Model Training
- 5) Evaluation
- 6) Prediction / Output



3) TEXT PREPROCESSING STEPS ☆

- Lowercasing
- Tokenization
- Stopword Removal
- Stemming / Lemmatization
- Removing Punctuation & Special Characters
- Handling Numbers
- Spelling Correction



4) TOKENIZATION ☆

Breaking text into smaller units called tokens.

- Word Tokenization → "NLP is fun" → ["NLP", "is", "fun"]
- Sentence Tokenization → Split into sentences
- Subword Tokenization → e.g. using BPE, WordPiece

5) FEATURE REPRESENTATION ☆

- Bag of Words (BoW)
- TF-IDF (Term Frequency - Inverse Document Frequency)
- Word Embeddings (Word2Vec, GloVe)
- Contextual Embeddings (BERT, RoBERTa, etc.)

6) POPULAR NLP MODELS ☆

- Naive Bayes
- Logistic Regression
- SVM (Support Vector Machine)
- RNN (Recurrent Neural Network)
- LSTM / GRU
- Transformers (BERT, GPT, T5, etc.)



7) EVALUATION METRICS ☆

- Accuracy
- Precision
- Recall
- F1-Score
- Confusion Matrix
- BLEU Score (for Translation)
- ROUGE Score (for Summarization)

8) NLP LIBRARIES IN PYTHON ☆

- NLTK (Natural Language Toolkit)
- spaCy
- TextBlob
- Gensim
- Hugging Face Transformers

9) APPLICATIONS OF NLP ☆

- Chatbots & Virtual Assistants
- Language Translation
- Sentiment Analysis
- Text Summarization
- Voice Recognition
- Spam Detection
- Search Engine
- Content Recommendation

10) CHALLENGES IN NLP ☆

- Ambiguity in Language
- Sarcasm & Irony
- Context Understanding
- Multilingual Support
- Data Quality
- Bias in Models



NLP is powering the future of communication between humans and machines.
The better we understand language, the smarter the machines become.



NATURAL LANGUAGE PROCESSING (NLP)




11) PARTS OF SPEECH (POS) TAGGING ☆

Assigns a part of speech to each word in a sentence.

Common Tags:

- Noun (NN)
- Verb (VB)
- Adjective (JJ)
- Adverb (RB)
- Pronoun (PRP)
- Preposition (IN)
- Conjunction (CC)
- Determiner (DT)
- Interjection (UH)



12) NAMED ENTITY RECOGNITION ☆

Identifies and classifies important entities in text into predefined categories.

Common Entities:

- Person (PER)
- Organization (ORG)
- Location (LOC)
- Date / Time (DATE)
- Money (MONEY)
- Percent (PERCENT)
- Product (PRODUCT)




13) SYNTACTIC ANALYSIS ☆

Analyzes the grammatical structure of a sentence.

Techniques:

- Parsing
- Constituency Parsing
- Dependency Parsing



14) SEMANTIC ANALYSIS ☆

Focuses on the meaning of the text.

Tasks:

- Word Sense Disambiguation (WSD)
- Semantic Role Labeling (SRL)
- Coreference Resolution
- Sentiment / Intent Understanding




15) WORD EMBEDDINGS ☆

Convert words into dense vector representations capturing meaning.

Popular Methods:

- Word2Vec (CBOW, Skip-gram)
- GloVe
- FastText
- BERT Embeddings



16) LANGUAGE MODELS ☆

Models that predict the next word in a sequence.

Types:

- N-gram Models
- Statistical Language Models
- Neural Language Models
- Pretrained LMs (GPT, BERT, LLaMA, etc.)



17) TEXT CLASSIFICATION ☆

Assigns a category/label to a given text.

Examples:

- Spam vs Not Spam
- Sentiment (Positive / Negative / Neutral)
- Topic Classification




18) SEQUENCE LABELING TASKS ☆

Assigns a label to each word in a sequence.

Examples:

- POS Tagging
- NER
- Chunking



19) MACHINE TRANSLATION ☆

Automatically translates text from one language to another.

Approaches:

- Rule-based MT
- Statistical MT (SMT)
- Neural MT (NMT - Seq2Seq, Transformer)



20) SENTIMENT ANALYSIS ☆

Determines the emotional tone behind the text.

Classes:



21) TOPIC MODELING ☆

Discovers hidden topics in a collection of documents.

Popular Algorithms:

- LDA (Latent Dirichlet Allocation)
- NMF (Non-negative Matrix Factorization)



22) CHATBOTS & DIALOG SYSTEMS ☆

Build conversational agents that can understand and respond to users.

Components:

- Intent Detection
- Entity Extraction
- Dialogue Management
- Response Generation



★ Understanding language is the first step towards building intelligent machines. 🚀

✧ NATURAL LANGUAGE PROCESSING (NLP) ✧




23) LANGUAGE MODELING ☆

Predicts the probability of a sequence of words.

Types:

- Unigram Model
- Bigram Model
- Trigram Model
- N-gram Model



24) PERPLEXITY ☆

Measures how well a probability model predicts a sample.

$$\text{Perplexity} = 2^{-\frac{1}{N} \sum_{i=1}^N \log_2 P(w_i)}$$


- Lower perplexity → Better model ✓
- Higher perplexity → Poorer model ✗

25) WORD SENSE DISAMBIGUATION (WSD) ☆

Finds the correct meaning of a word based on its context.

Approaches:

- Knowledge-based
- Supervised Learning
- Embedding-based



26) COREFERENCE RESOLUTION ☆

Identifies when different expressions refer to the same entity.

Example:

Riya went to the market. She bought some fruits.


Same Entity

27) PARSING ☆

Analyzes the grammatical structure of a sentence.

Types:

- Constituency Parsing (Tree structure)
- Dependency Parsing (Relations)




28) TRANSFORMERS ARCHITECTURE ☆

A deep learning model based on self-attention mechanism.

Key Components:

- Multi-Head Self-Attention
- Positional Encoding
- Encoder - Decoder Structure
- Feed Forward Networks



29) ATTENTION MECHANISM ☆

Allows the model to focus on relevant parts of the input.

Types:

- Scaled Dot-Product Attention
- Multi-Head Attention



30) Positional Encoding ☆


Adds position information to word embeddings since transformers have no recurrence.

Formula:

$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$


$$PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

Where,
pos = position, i = dimension index,
 d_{model} = embedding dimension




31) BERT (Bidirectional Encoder Representations from Transformers) ☆

- Pre-trained using Masked Language Modeling (MLM) and Next Sentence Prediction (NSP)
- Bidirectional context understanding
- Produces contextual embeddings



32) GPT (Generative Pre-trained Transformer) ☆

- Decoder-only architecture
- Trained using causal (autoregressive) language modeling
- Generates text one token at a time



33) TEXT GENERATION TECHNIQUES ☆

- Greedy Search → picks highest probability token
- Beam Search → keeps top k candidates
- Top-k Sampling → sample from top k tokens
- Top-p (Nucleus) Sampling → sample from top tokens with cumulative prob p
- Temperature Sampling → controls randomness



34) FINE-TUNING & TRANSFER LEARNING ☆

- Fine-tuning adapts a pre-trained model to a specific task with smaller dataset.
- Transfer learning saves time, data and improves performance.



★ NLP combines linguistics, statistics and deep learning to help machines understand human language. ♥



✦ NATURAL LANGUAGE PROCESSING (NLP) ✦



35) LANGUAGE UNDERSTANDING TASKS ☆

Help machines understand the meaning behind text.

Examples:

- Intent Classification
- Question Classification
- Natural Language Inference (NLI)
- Textual Entailment
- Semantic Similarity



36) SLOT FILLING ☆


Extracts specific information (data) from the user input based on predefined slots.

Example:

User Input	Extracted Slots
Book a flight from Hyderabad to Delhi on 25 May.	From : Hyderabad To : Delhi Date : 25 May

37) DIALOGUE STATE TRACKING ☆

Tracks the context and state of a conversation over multiple turns.



38) TEXT NORMALIZATION ☆

Converts informal or noisy text into a standard format.

Examples:

- "gud mrng" → "good morning"
- "12th Jan '24" → "12 January 2024"
- "can't" → "cannot"



39) LEMMATIZATION VS STEMMING ☆

Stemming	Lemmatization
• Removes word endings • May not be real words • Simple & Fast	• Converts to root word (dictionary form) • Always real words • Slower but accurate


Examples:

Word	Stemming	Lemmatization
running	run	run
better	better	good
studies	studi	study

40) STOP WORDS ☆

Common words in a language that do not carry much meaning and are often removed.

Examples:



41) N-GRAMS ☆

A sequence of N items (words/characters) from a given text.

Types:

- Unigram (1-gram): one word e.g., I love NLP
- Bigram (2-gram): two words e.g., I love, love NLP
- Trigram (3-gram): three words e.g., I love NLP



42) HANDLING OUT-OF-VOCABULARY (OOV) ☆

Words that are not present in the vocabulary.

Solutions:

- Use <UNK> token
- Subword Tokenization (BPE, WordPiece)
- Character-level models



43) SUBWORD TOKENIZATION ☆

Breaks words into smaller units (subwords) to handle rare or unseen words.

Methods:

- Byte Pair Encoding (BPE)
- WordPiece
- SentencePiece

Example (BPE):

unbelievable → un + belief + able

44) LANGUAGE IDENTIFICATION ☆

Detects the language of a given text.

Example:



45) CODE-MIXING & CODE-SWITCHING ☆

Mixing multiple languages in a single sentence or switching between languages.

Examples:

- Code-Mixing: "Kal meeting hai, ready ho na?" (Hindi + English)
- Code-Switching: "I will call you later, ठीक है?" (Switching between English & Hindi)



46) PRAGMATICS IN NLP ☆

Understanding language in context, considering speaker's intent, tone, and real-world knowledge.

Example:

- "Can you pass the salt?" → Not a question, but a request.



NLP is not just about processing text, it's about understanding humans.
The more context we give machines, the better they understand us. ♥



teaching4u



MULTIMODAL NLP (M-NLP)



Multimodal NLP enables models to understand and generate responses using multiple types of data (modalities) like text, images, audio, and video together.

1) WHAT IS MULTIMODAL NLP? ☆

Traditional NLP works only with text. Multimodal NLP combines two or more modalities to understand context better and generate more accurate and human-like responses.



2) COMMON MODALITIES ☆

- Text – words, sentences, documents
- Image – photos, diagrams, charts
- Audio – speech, sounds
- Video – frames, actions, scenes
- Sensor Data – GPS, IoT, signals, etc.




Combining these helps models understand richer context!

3) WHY MULTIMODAL NLP? ☆

- Closer to human communication
- Better understanding of context
- Handles real-world complex data
- Improves accuracy & robustness
- Enables new applications



4) MULTIMODAL NLP ARCHITECTURE ☆



Encoders convert each modality → Fusion layer combines → LLM generates final response

5) EXAMPLES OF MULTIMODAL NLP ☆

- Image + Text → Visual Question Answering (e.g., What is in this picture?)
- Audio + Text → Speech to Text + Understanding (e.g., Transcribe and summarize)
- Video + Text → Video Captioning / Q&A (e.g., Describe this video)
- Image + Text + Audio → Multimodal Chatbots / Assistants.

Used in: Chatbots, Content Analysis, Accessibility, Healthcare, Education, Search, Autonomous Systems & more. ☆

6) POPULAR MULTIMODAL MODELS ☆

- CLIP (Text + Image)
- BLIP / BLIP-2 (Image-Text Understanding & Generation)
- Flamingo (DeepMind)
- LLaVA (Large Language and Vision Assistant)
- PaLI (Google)
- GPT-4V / Gemini (Vision + Text)
- Whisper (Audio + Text)
- Video-LLaVA (Video + Text)



7) CHALLENGES ☆

- Data alignment across modalities
- High computational cost
- Handling missing or noisy modalities
- Bias & fairness in multimodal data
- Long context understanding



8) FUTURE OF MULTIMODAL NLP ☆

- More human-like AI interactions
- Real-time multimodal assistants
- Better accessibility & inclusivity
- Strong impact on Robotics, AR/VR, Healthcare, Education & beyond!



Multimodal NLP is the bridge between how machines perceive and how humans communicate! ❤️